

Automating Keyboard Typing


Practical Reinforcement Learning



Aidan Keaveny

Overview

1. Problem Formulation
2. What is Reinforcement Learning?
3. Practical RL for URC Competition
4. Future Work

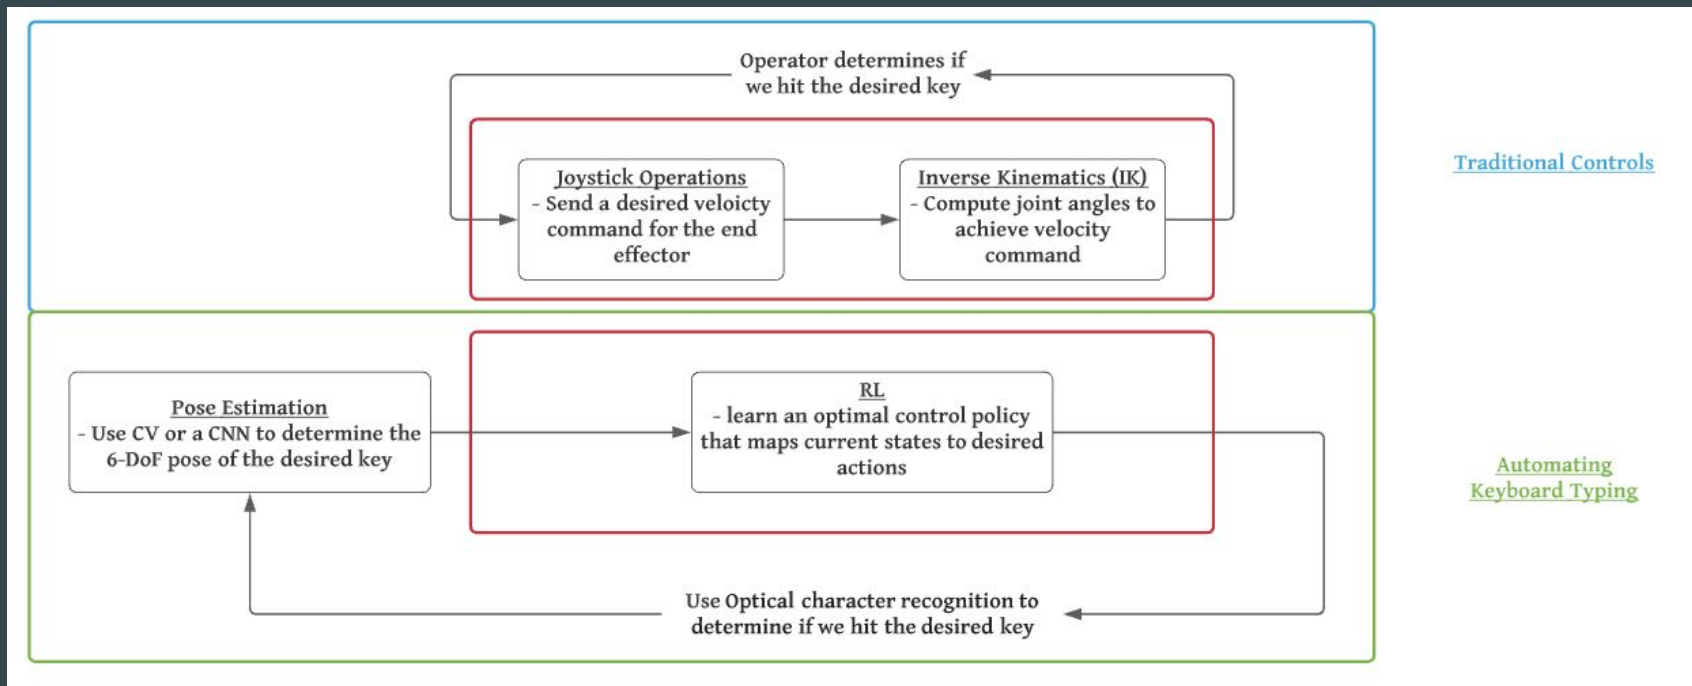


Problem Formulation

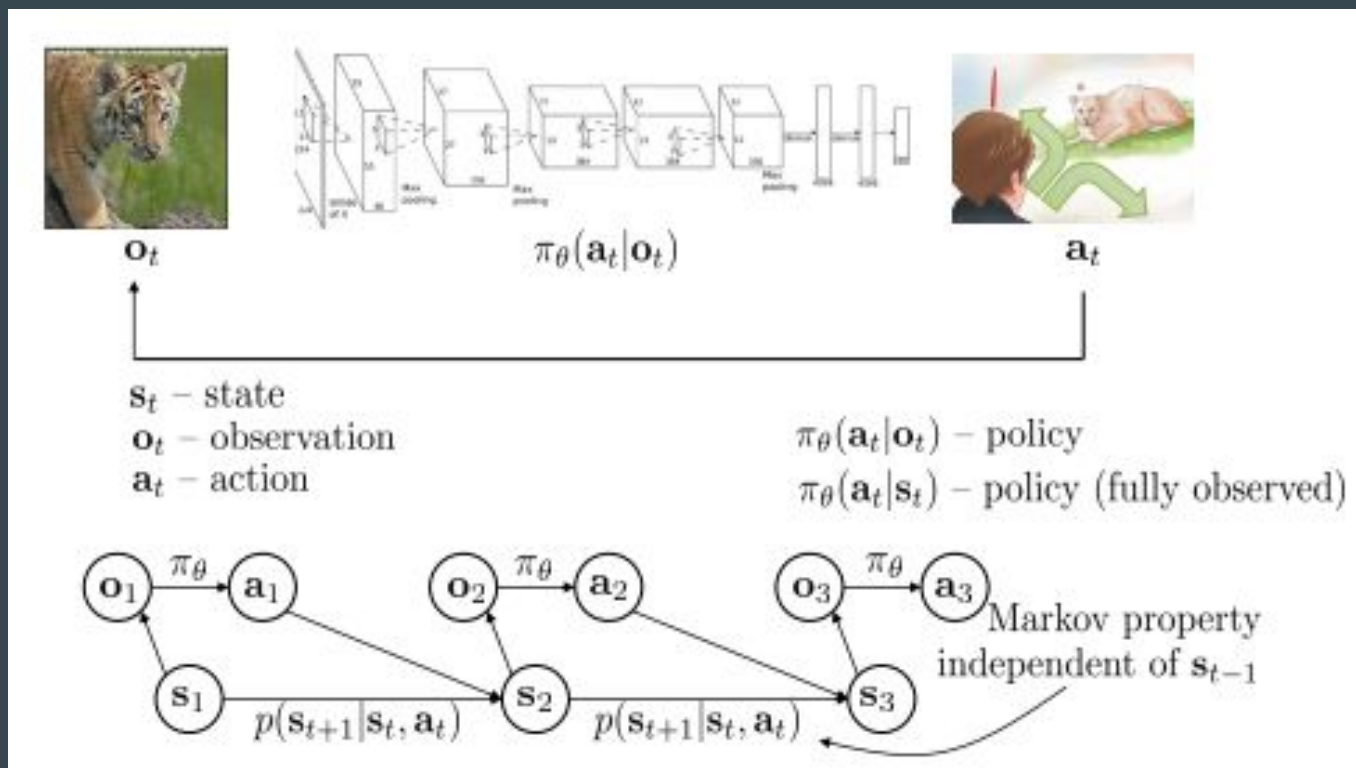
- Keyboard typing is a *high precision task* at URC competition
- Main Goal: Attempt to replace traditional path planning techniques or joystick operations with an optimal control policy that is learnt via Reinforcement Learning.



Problem Formulation

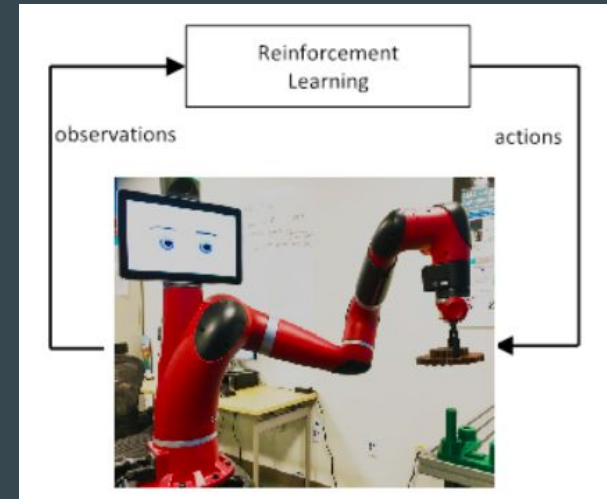


What is RL?



What is RL?


- Env: The agent is always acting in an environment.
- State-Actions Pairs: The agent in one of many states (s) of the environment can choose to take one of many actions (a).
- Model: How the environment reacts to certain actions is defined by a model which we may or may not know.
- Reward: Once an action is taken, the environment delivers a reward (r) as feedback.



What is RL?



- Goal: We want to hit a desired key so we need the 6-DoF Pose ...
- States s_t [1x39]: goal, initial pose, *current* joint positions, velocities, efforts, position and velocity limit switches
- Actions a_t [1x5]: *desired* joint positions (position control)
- Reward Function R :

$$\% \text{ distance remaining} = \frac{|\text{distance to target}|}{|\text{initial distance}|}$$



RL Overview

- Game playing: AlphaGo involves both model-free methods (CNN), and also model-based methods (Monte Carlo Tree Search)
- Control problems: much less structure..
 - Effective representations of the state space s_t
 - Discrete or continuous control actions a_t
 - Engineered reward functions R



RL Overview



Atari games:

Q-learning:

V. Mnih, K. Kavukcuoglu, D. Silver, A. Graves, I. Antonoglou, et al. "Playing Atari with Deep Reinforcement Learning". (2013).

Policy gradients:

J. Schulman, S. Levine, P. Moritz, M. I. Jordan, and P. Abbeel. "Trust Region Policy Optimization". (2015).
V. Mnih, A. P. Badia, M. Mirza, A. Graves, T. P. Lillicrap, et al. "Asynchronous methods for deep reinforcement learning". (2016).



Real-world robots:

Guided policy search:

S. Levine*, C. Finn*, T. Darrell, P. Abbeel. "End-to-end training of deep visuomotor policies". (2015).

Q-learning:

D. Kalashnikov et al. "QT-Opt: Scalable Deep Reinforcement Learning for Vision-Based Robotic Manipulation". (2018).

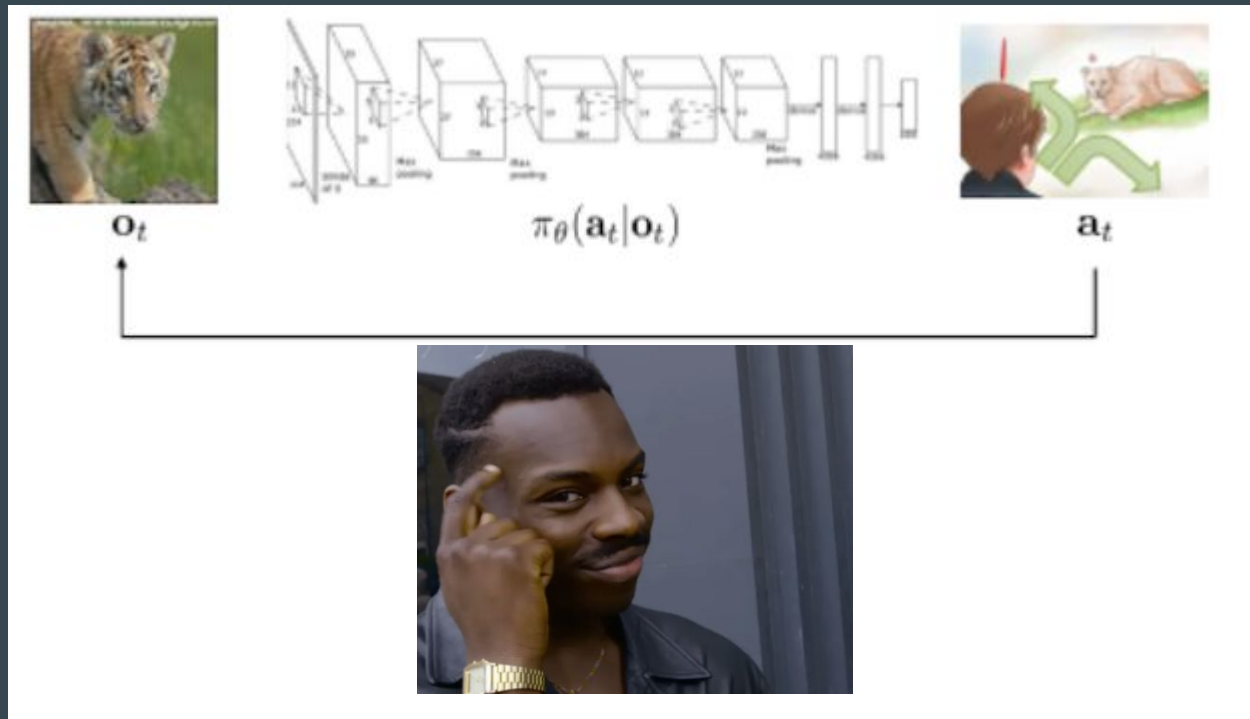


Beating Go champions:


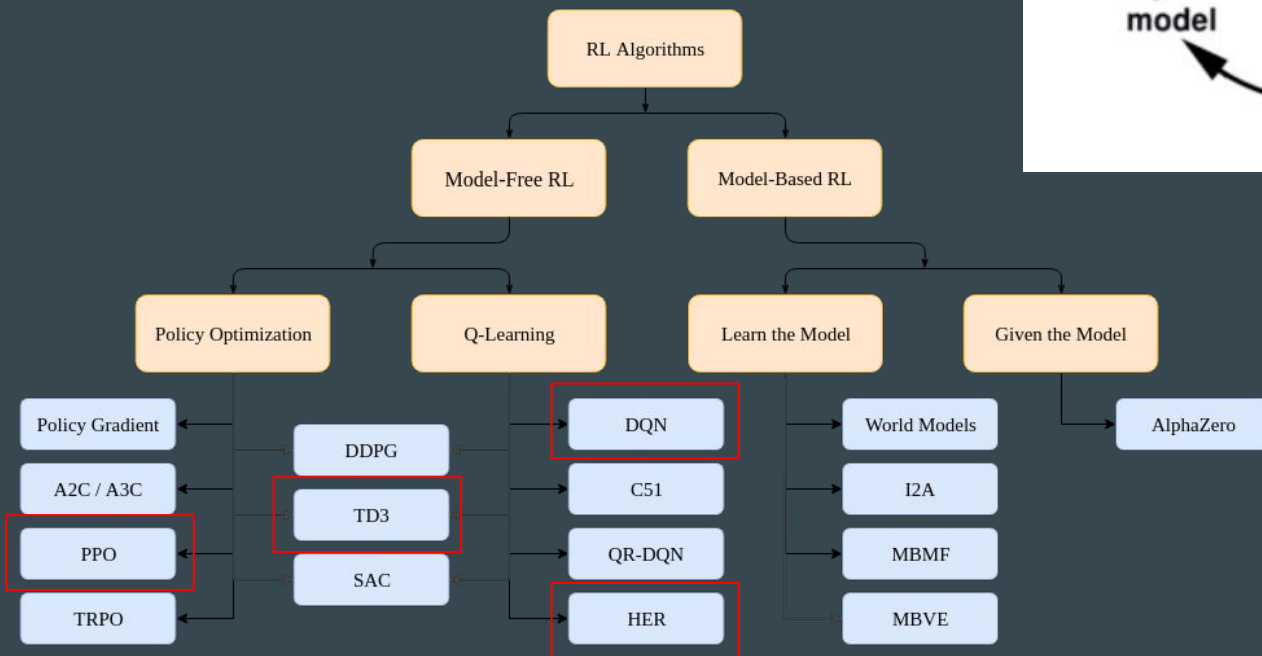
Supervised learning + policy
gradients + value functions +
Monte Carlo tree search:

D. Silver, A. Huang, C. J. Maddison, A. Guez, L. Sifre, et al. "Mastering the game of Go with deep neural networks and tree search". Nature (2016).

What is RL?



Different Types of RL




Different Types of RL

Types of RL algorithms

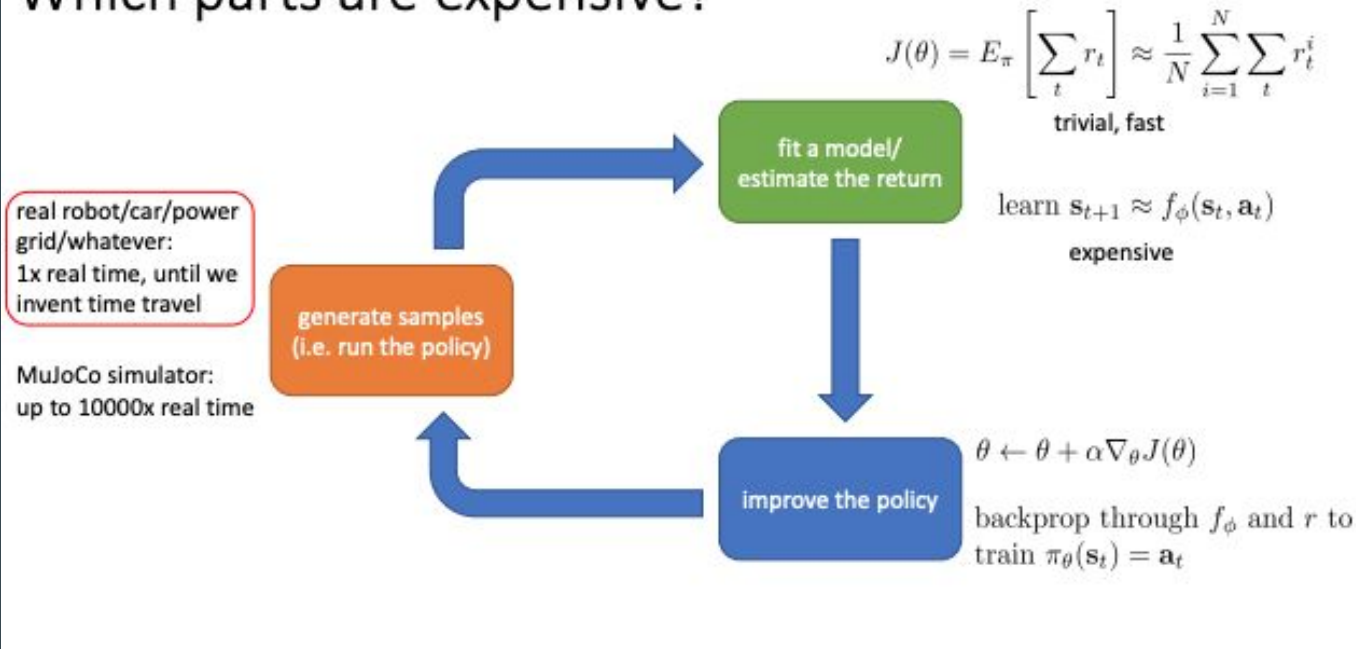
$$\theta^* = \arg \max_{\theta} E_{\tau \sim p_{\theta}(\tau)} \left[\sum_t r(s_t, a_t) \right]$$

- Policy gradients: directly differentiate the above objective
- Value-based: estimate value function or Q-function of the optimal policy (no explicit policy)
- Actor-critic: estimate value function or Q-function of the current policy, use it to improve policy
- Model-based RL: estimate the transition model, and then...
 - Use it for planning (no explicit policy)
 - Use it to improve a policy
 - Something else



RL Overview


Which parts are expensive?



RL Overview

Why so many RL algorithms?


- Different tradeoffs
 - Sample efficiency
 - Stability & ease of use
- Different assumptions
 - Stochastic or deterministic?
 - Continuous or discrete?
 - Episodic or infinite horizon?
- Different things are easy or hard in different settings
 - Easier to represent the policy?
 - Easier to represent the model?



Different Types of RL

- Learning from demonstrations
 - Directly copying observed behavior
 - Inferring rewards from observed behavior (inverse reinforcement learning)
- Learning from observing the world
 - Learning to predict *"yeah, keyboard typing isn't hard enough .."*
 - Unsupervised learning
- Learning from other tasks
 - Transfer learning
 - Meta-learning: learning to learn

Small break before Practical RL ..



OpenAI + Stable-Baselines3

```
class TestClass:
    NUM_EPISODES = 1
    MAX_STEPS = 5000
    KEY_POSITION = np.array([0.6, 0.6, 0.6])
    KEY_ORIENTATION = np.array([0, 0, 0, 1])



    def __run_test(self, env):
        pp = pprint.PrettyPrinter() # TODO: update to python 3.8 to use sort_dicts = False

        for episode in range(self.NUM_EPISODES):
            initial_observation = env.reset()
            print('Initial Observation:')
            pp.pprint(initial_observation)

            for sim_step in range(self.MAX_STEPS):
                action = env.action_space.sample()
                observation, reward, done, info = env.step(action)

                print()
                print('Action:')
                pp.pprint(action)
                print('Observation:')
                pp.pprint(observation)
                print('Info:')
                pp.pprint(info)
                print('Reward:')
                pp.pprint(reward)

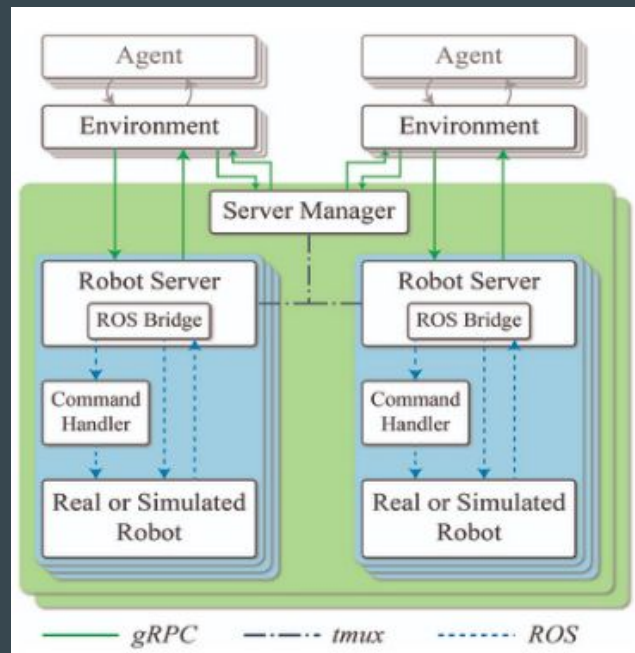

            if done:
                print()
                print(f'Episode #{episode} finished after {info["sim"]["steps_executed"]} steps!')
                print(f'Episode #{episode} exit condition was {info["sim"]["end_condition"]}')
                print()
                break
```



Progress Update: SAR





Robo-Gym: *Sim 2 Real* ..




GAZEBO



Pose Estimation: *my thesis work..*

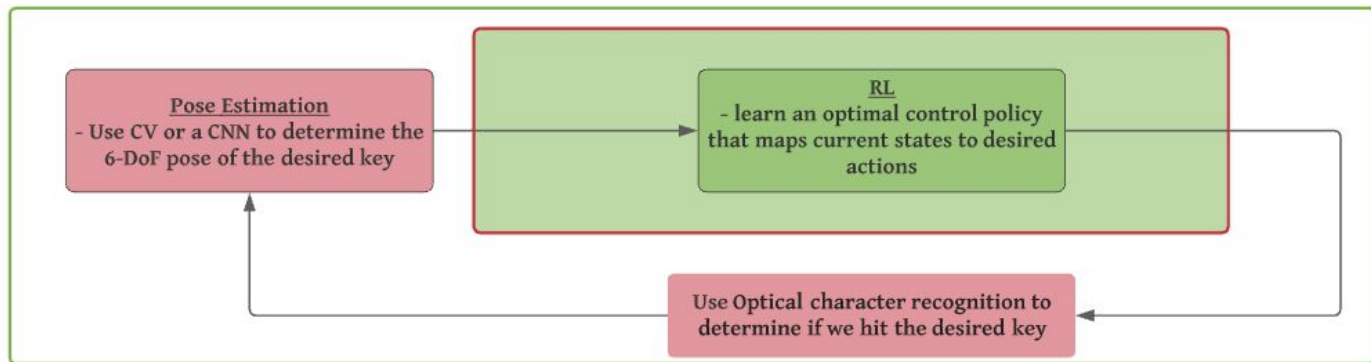


Robohub's Kinova Gen3



Future Work

- Arm: Test arm's IK with closed loop feedback
- RL: Test RL Algorithm with real hardware
- Pose Estimation: Experiments with fixed keyboard position ...
- Robohub: Can always use the Gen3 arm for testing



Automating
Keyboard Typing

References

- [AlphaGo Documentary](#)
- [Sutton & Barto Textbook \(U of A !!!\)](#)
- [CS285: Berkeley's Deep Reinforcement Learning](#)
- [OpenAI Gym Tutorial](#)
- [Stable -Baselines3](#)
- [Robo-Gym](#)